

Data Transformation and Exploration

Lecture 03.2: Relational Data

Lauren Sullivan

Module: Data Management, Visualization & Reproducibility

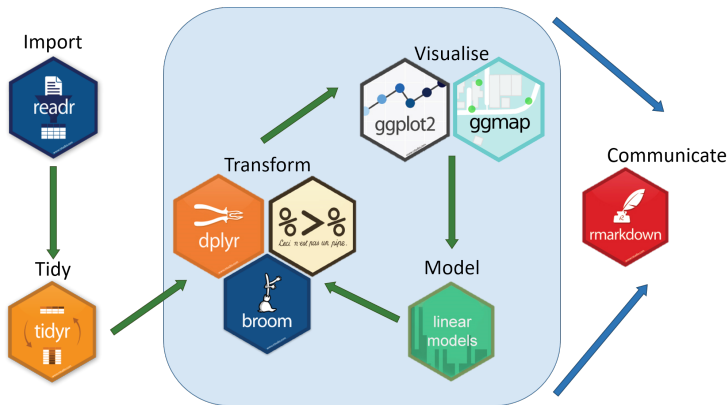
Relational Data

As we mentioned previously, it is good practice to maintain smaller datasets and then merge them together through code. Using multiple tables of data is called *Relational Data* because we are interested in the relations between datasets, not individual ones.

Relations are always built between pairs of tables. And to do this work we need some terminology.

- ▶ **Mutating joins:** adds a new variable to one data frame from matching observations in another
- ▶ **Filtering joins:** filters observations from one data frame based on if they match an observation from another data frame.
- ▶ **Set operations:** treats observations as if they were set elements.

Relational Data with dplyr



In this lecture we will use the `library(nycflights13)` from [R for Data Science](#). This library contains data on flights in and out of NYC.

```
library(nycflights13)
```


```
airlines  # describes all airlines in NYC
```

```
airports  # describes the airports flights go to/from
```

```
planes    # describes each plane
```

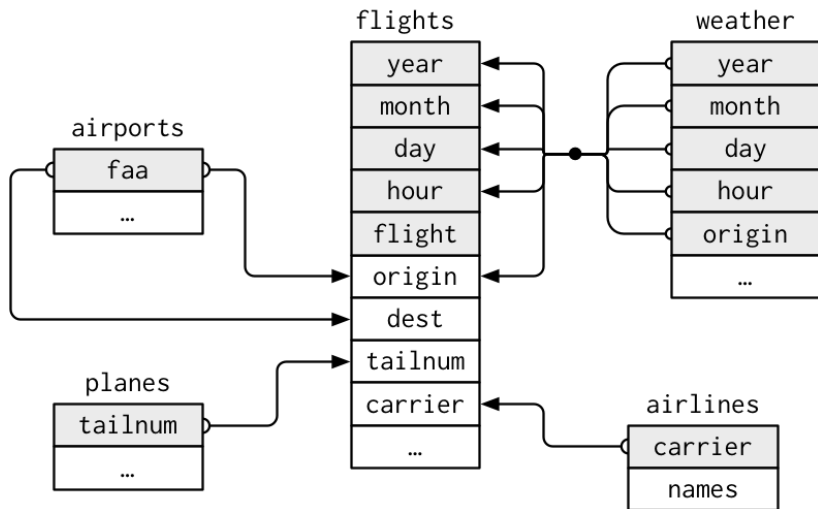
```
weather   # weather at each NYC airport each hour
```

For More Information




R for Data Science.

Relational data for nycflights13



Relational data for nycflights13



- ▶ **flights** connects to **planes** via a single variable, `tailnum`.
- ▶ **flights** connects to **airlines** through the `carrier` variable.
- ▶ **flights** connects to **airports** in two ways: via the `origin` and `dest` variables.
- ▶ **flights** connects to **weather** via `origin` (the location), and `year`, `month`, `day` and `hour` (the time).

Keys

The terminology used by *R for Data Science* for a variable (or set of variables) that connects each pair of tables.

- ▶ There are three types of keys
 1. *Primary Key* - uniquely identifies an observation in its own table. Here, `planes$tailnum` uniquely identifies each plane in the `planes` table.
 2. *Foreign Key* - uniquely identifies an observation in another table. Here, `flights$tailnum` appears in the `flights` table where it matches each flight to a unique plane in `planes`.
 3. *Surrogate Key* - sometimes datasets do not contain a primary key, so you must make one. Here, `flights$tailnum` is not a primary key because the same plane appears multiple times per day at the airport. Create one using, for example, `mutate()` and `row_number()`. **This then becomes a *primary key*.**

A *primary key* that corresponds to a *foreign key* is a relation. This builds the 1:1, 1:n, or n:1 relationships we mentioned previously.

Mutating Joins

Mutating joins combine variables from two tables.

Let's make `flights` smaller so we can better see our data.

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)  
flights2
```

```
## # A tibble: 336,776 x 8  
##   year month   day hour origin dest  tailnum carrier  
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr>  
## 1  2013     1     1     5 EWR   IAH   N14228  UA  
## 2  2013     1     1     5 LGA   IAH   N24211  UA  
## 3  2013     1     1     5 JFK   MIA   N619AA  AA  
## 4  2013     1     1     5 JFK   BQN   N804JB  B6  
## 5  2013     1     1     6 LGA   ATL   N668DN  DL  
## 6  2013     1     1     5 EWR   ORD   N39463  UA  
## 7  2013     1     1     6 EWR   FLL   N516JB  B6  
## 8  2013     1     1     6 LGA   IAD   N829AS  EV  
## 9  2013     1     1     6 JFK   MCO   N593JB  B6  
## 10 2013     1     1     6 LGA   ORD   N3ALAA  AA  
## # ... with 336,766 more rows
```

Mutating Joins

To add the full name of the airline to `flights2`, you combine `airlines` with `flights2` with `left_join()`.

```
flights2 %>%  
  select(-origin, -dest) %>%  
  left_join(airlines, by = "carrier")
```


```
## # A tibble: 336,776 x 7
```

```
##   year month   day hour tailnum carrier name  
##   <int> <int> <int> <dbl> <chr>   <chr>   <chr>  
## 1  2013     1     1     5 N14228   UA      United Air Lines Inc.  
## 2  2013     1     1     5 N24211   UA      United Air Lines Inc.  
## 3  2013     1     1     5 N619AA   AA      American Airlines Inc.  
## 4  2013     1     1     5 N804JB   B6      JetBlue Airways  
## 5  2013     1     1     6 N668DN   DL      Delta Air Lines Inc.  
## 6  2013     1     1     5 N39463   UA      United Air Lines Inc.  
## 7  2013     1     1     6 N516JB   B6      JetBlue Airways  
## 8  2013     1     1     6 N829AS   EV      ExpressJet Airlines Inc.  
## 9  2013     1     1     6 N593JB   B6      JetBlue Airways  
## 10 2013     1     1     6 N3ALAA   AA      American Airlines Inc.
```


```
## # ... with 336,766 more rows
```

Joins, How Do They Work?

Imagine you have two datasets, x (on the left) and y (on the right), and they each have *key* variables (colored column). The following diagram shows all potential matches with the grid lines




However the only possible matches are indicated with colored dots. The resulting dataset demonstrates that in this case all variables are joined, and some observations are lost.



Inner Joins

Inner joins match pairs of observations based on equal key variables, but they lose observations without matches.

- ▶ You specify the key variable with `by`.



```
x %>%  
  inner_join(y, by = "key")
```

```
## # A tibble: 2 x 3  
##   key val_x val_y  
##   <dbl> <chr> <chr>  
## 1     1 x1    y1  
## 2     2 x2    y2
```




Outer Joins

Outer joins also match observations based on key variables, but **do not** lose observations.

- ▶ There are three types of outer joins.
 1. A **left join** keeps all observations in **x**.
 2. A **right join** keeps all observations in **y**.
 3. A **full join** keeps all observations in **x** and **y**.

In all cases, unmatched observations are filled with **NA**.


Outer Joins



Duplicate Keys

Sometimes keys are not unique. There are two possibilities here.

1. One table has duplicate keys (a 1:n relationship)




```
left_join(x, y, by = "key")
```

```
## # A tibble: 4 x 3
##   key val_x val_y
##   <dbl> <chr> <chr>
## 1     1  x1    y1
## 2     2  x2    y2
## 3     2  x3    y2
## 4     1  x4    y1
```

Duplicate Keys

- Both tables have duplicate keys. This is often an error because neither table has a unique identifier.



```
left_join(x, y, by = "key")
```

```
## # A tibble: 6 x 3
##   key val_x val_y
##   <dbl> <chr> <chr>
## 1     1  x1    y1
## 2     2  x2    y2
## 3     2  x2    y3
## 4     2  x3    y2
## 5     2  x3    y3
## 6     3  x4    y4
```


Defining the Key Columns

Natural joins, or `by = NULL` uses all variables that are in both tables. The `flights` and `weather` table have `year`, `month`, `day`, `hour`, and `origin` in common.

```
flights2 %>%  
  left_join(weather)
```

```
## Joining, by = c("year", "month", "day", "hour", "origin")
```

```
## # A tibble: 336,776 x 18
```

```
##   year month   day hour origin dest tailnum carrier temp dewp humid  
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>  
## 1 2013     1     1     5 EWR  IAH  N14228  UA      39.0  28.0  64.4  
## 2 2013     1     1     5 LGA  IAH  N24211  UA      39.9  25.0  54.8  
## 3 2013     1     1     5 JFK  MIA  N619AA  AA      39.0  27.0  61.6  
## 4 2013     1     1     5 JFK  BQN  N804JB  B6      39.0  27.0  61.6  
## 5 2013     1     1     6 LGA  ATL  N668DN  DL      39.9  25.0  54.8  
## 6 2013     1     1     5 EWR  ORD  N39463  UA      39.0  28.0  64.4  
## 7 2013     1     1     6 EWR  FLL  N516JB  B6      37.9  28.0  67.2  
## 8 2013     1     1     6 LGA  IAD  N829AS  EV      39.9  25.0  54.8  
## 9 2013     1     1     6 JFK  MCO  N593JB  B6      37.9  27.0  64.3  
## 10 2013     1     1     6 LGA  ORD  N3ALAA  AA      39.9  25.0  54.8  
## # ... with 336,766 more rows, and 7 more variables: wind_dir <dbl>,  
## #   wind_speed <dbl>, wind_gust <dbl>, precip <dbl>, pressure <dbl>,  
## #   visib <dbl>, time_hour <dtm>
```

Defining the Key Columns

You can join by a *character vector* with `by = "x"`. Try joining `flights` and `planes` by `tailnum`.

```
flights2 %>%  
  left_join(planes, by = "tailnum")
```

```
## # A tibble: 336,776 x 16  
##   year.x month   day   hour origin dest  tailnum carrier year.y type    manuf-1  
##   <int> <int> <int> <dbl> <chr>  <chr> <chr>  <chr>    <int> <chr>  <chr>  
## 1  2013     1     1     5   EWR   IAH   N14228  UA      1999 Fixed w~ BOEING  
## 2  2013     1     1     5   LGA   IAH   N24211  UA      1998 Fixed w~ BOEING  
## 3  2013     1     1     5   JFK   MIA   N619AA  AA      1990 Fixed w~ BOEING  
## 4  2013     1     1     5   JFK   BQN   N804JB  B6      2012 Fixed w~ AIRBUS  
## 5  2013     1     1     6   LGA   ATL   N668DN  DL      1991 Fixed w~ BOEING  
## 6  2013     1     1     5   EWR   ORD   N39463  UA      2012 Fixed w~ BOEING  
## 7  2013     1     1     6   EWR   FLL   N516JB  B6      2000 Fixed w~ AIRBUS~  
## 8  2013     1     1     6   LGA   IAD   N829AS  EV      1998 Fixed w~ CANADA~  
## 9  2013     1     1     6   JFK   MCO   N593JB  B6      2004 Fixed w~ AIRBUS  
## 10 2013     1     1     6   LGA   ORD   N3ALAA  AA      NA <NA>    <NA>  
## # ... with 336,766 more rows, 5 more variables: model <chr>, engines <int>,  
## #   seats <int>, speed <int>, engine <chr>, and abbreviated variable name  
## #   1: manufacturer
```

You can see that both datasets had a `year` column but they had different values, so the resulting variables are `year.x` and `year.y`.

Defining the Key Columns

A *named character vector*, for example `by = c("a" = "b")` allows you to match variable `a` in table `x` to variable `b` in table `y`. Note: the variables from `x` will always be used in the output.


```
flights2 %>%  
  left_join(airports, c("dest" = "faa"))
```

```
## # A tibble: 336,776 x 15  
##   year month   day hour origin dest tailnum carrier name   lat lon alt  
##   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <chr> <dbl> <dbl> <dbl>  
## 1 2013     1     1     5 EWR   IAH  N14228  UA    Georg~ 30.0 -95.3  97  
## 2 2013     1     1     5 LGA   IAH  N24211  UA    Georg~ 30.0 -95.3  97  
## 3 2013     1     1     5 JFK   MIA  N619AA  AA    Miami~ 25.8 -80.3   8  
## 4 2013     1     1     5 JFK   BQN  N804JB  B6    <NA>   NA    NA    NA  
## 5 2013     1     1     6 LGA   ATL  N668DN  DL    Harts~ 33.6 -84.4 1026  
## 6 2013     1     1     5 EWR   ORD  N39463  UA    Chica~ 42.0 -87.9  668  
## 7 2013     1     1     6 EWR   FLL  N516JB  B6    Fort ~ 26.1 -80.2   9  
## 8 2013     1     1     6 LGA   IAD  N829AS  EV    Washi~ 38.9 -77.5  313  
## 9 2013     1     1     6 JFK   MCO  N593JB  B6    Orlan~ 28.4 -81.3  96  
## 10 2013     1     1     6 LGA   ORD  N3ALAA  AA    Chica~ 42.0 -87.9  668  
## # ... with 336,766 more rows, and 3 more variables: tz <dbl>, dst <chr>,  
## #   tzone <chr>
```


Filtering Joins

Filtering joins only affect observations, not variables.

- **semi_join(x, y)** **keeps** observations in **x** with a match in **y**.



- **anti_join(x, y)** **drops** observations in **x** with a match in **y**.



semi_join()

So say you have a table of the top destinations (`top_dest`) and you want to know what flights go to those destinations.

```
trips <- flights %>%  
  semi_join(top_dest)
```

```
## Joining, by = "dest"
```

```
trips[,1:7]
```

```
## # A tibble: 141,145 x 7
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
## 1  2013     1     1     542           540           2     923  
## 2  2013     1     1     554           600          -6     812  
## 3  2013     1     1     554           558          -4     740  
## 4  2013     1     1     555           600          -5     913  
## 5  2013     1     1     557           600          -3     838  
## 6  2013     1     1     558           600          -2     753  
## 7  2013     1     1     558           600          -2     924  
## 8  2013     1     1     558           600          -2     923  
## 9  2013     1     1     559           559           0     702  
## 10 2013     1     1     600           600           0     851  
## # ... with 141,135 more rows
```

anti_join()

Anti-joins are useful for figuring out mismatches. So say you want to know the `flights` that do not have a match in `planes`.

```
mismatch <- flights %>%  
  anti_join(planes, by = "tailnum")  
mismatch[,1:7]
```

```
## # A tibble: 52,606 x 7  
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>  
## 1  2013     1     1     558             600        -2     753  
## 2  2013     1     1     559             600        -1     941  
## 3  2013     1     1     600             600         0     837  
## 4  2013     1     1     602             605        -3     821  
## 5  2013     1     1     608             600         8     807  
## 6  2013     1     1     611             600        11     945  
## 7  2013     1     1     623             610        13     920  
## 8  2013     1     1     624             630        -6     840  
## 9  2013     1     1     628             630        -2    1137  
## 10 2013     1     1     629             630        -1     824  
## # ... with 52,596 more rows
```

Set Operations

Set operations work with complete rows, and compare the values of every variable. For set operations, it is assumed that **x** and **y** have the same variables, and treat all observations as sets.

- ▶ `intersect(x, y)` returns only observations in both **x** and **y**.
- ▶ `union(x, y)` returns unique observations in **x** and **y**.
- ▶ `setdiff(x, y)` returns observations in **x**, but not **y**.

Say you have these datasets.

```
df1 <- tribble(
  ~x, ~y,
  1,  1,
  2,  1
)
df2 <- tribble(
  ~x, ~y,
  1,  1,
  1,  2
)
```

Set Operations

```
intersect(df1, df2)
```

```
## # A tibble: 1 x 2
##       x     y
##   <dbl> <dbl>
## 1     1     1
```

```
union(df1, df2) #note we get 3 rows, not 4.
```

```
## # A tibble: 3 x 2
##       x     y
##   <dbl> <dbl>
## 1     1     1
## 2     2     1
## 3     1     2
```

```
setdiff(df1, df2)
```

```
## # A tibble: 1 x 2
##       x     y
##   <dbl> <dbl>
## 1     2     1
```

```
setdiff(df2, df1)
```

```
## # A tibble: 1 x 2
##       x     y
##   <dbl> <dbl>
## 1     1     2
```